

MA679 Hw4

Jiahao Xu


2/13/2019

5.8 (a)

```
set.seed(1)
y <- rnorm(100)
x <- rnorm(100)
y <- x - 2 * x^2 + rnorm(100)
# n is 100 and p is 2, model is  $Y=x-2x^2+error$ 
```

(b)

```
plot(x, y)
```



#(c)

```
library(boot)
set.seed(100)
data<-data.frame(x,y)
mod1<-glm(y ~ x)
cv1=cv.glm(data,mod1)$delta[1]
output=paste("When X is in poly degree 1, CV is", cv1)
output
```

```
## [1] "When X is in poly degree 1, CV is 5.89097855988842"
```

```

mod2<-glm(y~poly(x,2))
cv2=cv.glm(data,mod2)$delta[1]
output=paste("When X is in poly degree 2, CV is", cv2)
output

## [1] "When X is in poly degree 2, CV is 1.0865955642745"

mod3<-glm(y~poly(x,3))
cv3=cv.glm(data,mod3)$delta[1]
output=paste("When X is in poly degree 3, CV is", cv3)
output

## [1] "When X is in poly degree 3, CV is 1.10258509387339"

mod4<-glm(y~poly(x,4))
cv4=cv.glm(data,mod4)$delta[1]
output=paste("When X is in poly degree 4, CV is", cv4)
output

## [1] "When X is in poly degree 4, CV is 1.11477226814508"

```

(d)

```

set.seed(50)
data<-data.frame(x,y)
mod1<-glm(y ~ x)
cv1=cv.glm(data,mod1)$delta[1]
output=paste("When X is in poly degree 1, CV is", cv1)
output

## [1] "When X is in poly degree 1, CV is 5.89097855988843"

mod2<-glm(y~poly(x,2))
cv2=cv.glm(data,mod2)$delta[1]
output=paste("When X is in poly degree 2, CV is", cv2)
output

## [1] "When X is in poly degree 2, CV is 1.0865955642745"

mod3<-glm(y~poly(x,3))
cv3=cv.glm(data,mod3)$delta[1]
output=paste("When X is in poly degree 3, CV is", cv3)
output

## [1] "When X is in poly degree 3, CV is 1.10258509387339"

mod4<-glm(y~poly(x,4))
cv4=cv.glm(data,mod4)$delta[1]
output=paste("When X is in poly degree 4, CV is", cv4)
output

## [1] "When X is in poly degree 4, CV is 1.11477226814507"

```

The results from c and d are exactly the same because LOOCV evaluates n folds of a single observation

(e)

*#From the CV result, we can see that mod2 has the smallest value, which x is in 2 degree poly. It is wh
#I expect because in the part (a), we can see that the relation is quadratic.*

(f)

```
summary(mod4)
```

```
##
## Call:
## glm(formula = y ~ poly(x, 4))
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8914  -0.5244   0.0749   0.5932   2.7796
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -1.8277     0.1041  -17.549  <2e-16 ***
## poly(x, 4)1    2.3164     1.0415    2.224  0.0285 *
## poly(x, 4)2  -21.0586     1.0415  -20.220  <2e-16 ***
## poly(x, 4)3   -0.3048     1.0415   -0.293  0.7704
## poly(x, 4)4   -0.4926     1.0415   -0.473  0.6373
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 1.084654)
##
##      Null deviance: 552.21  on 99  degrees of freedom
## Residual deviance: 103.04  on 95  degrees of freedom
## AIC: 298.78
##
## Number of Fisher Scoring iterations: 2
```

*# From the summary, we can obviously realize that only intercept and quadratic term have the significant
value. This observation is the same with the result from crossvalidation.*

6.2

*# (a) Lasso is less flexible compared to linear regression since it has more restrictions
and will give improved prediction accuracy when its increase in bias less than its decrease in variance
(b) Ridge regression is less flexible compared to linear regression since it has more restrictions
and will give improved prediction accuracy when its increase in bias less than its decrease in variance
(c) Non-linear regression is more flexible compared to linear regression since it has no restrictions
and will give improved prediction accuracy when its increase in variance less than its decrease in bias*

6.10 (a)


```
set.seed(100)
x <- matrix(rnorm(1000 * 20), 1000, 20)
b <- rnorm(20)
b[1] <- 0
b[4] <- 0
b[3] <- 0
b[7] <- 0
b[19] <- 0
b[5] <- -0
error <- rnorm(1000)
y <- x %*% b + error
```

(b)

```
train <- sample(seq(1000), 100, replace = FALSE)
x.train <- x[train, ]
x.test <- x[-train, ]
y.train <- y[train]
y.test <- y[-train]
```

(c)


```
library(leaps)
traindata <- data.frame(y = y.train, x = x.train)
regfit.full <- regsubsets(y ~ ., data = traindata, nvmax = 20)
train.mat <- model.matrix(y ~ ., data = traindata, nvmax = 20)
val.errors <- rep(NA, 20)
for (i in 1:20) {
  coefi <- coef(regfit.full, id = i)
  pred <- train.mat[, names(coefi)] %*% coefi
  val.errors[i] <- mean((pred - y.train)^2)
}
plot(val.errors, xlab = "Number of predictors", ylab = "Training MSE", pch = 19, type = "b")
```



Number of predictors

#(d)

```
testdata <- data.frame(y = y.test, x = x.train)
regfit.full2 <- regsubsets(y ~ ., data = testdata, nvmax = 20)
test.mat <- model.matrix(y ~ ., data = testdata, nvmax = 20)
val.errors2 <- rep(NA, 20)
for (i in 1:20) {
  coefi <- coef(regfit.full2, id = i)
  pred <- test.mat[, names(coefi)] %*% coefi
  val.errors2[i] <- mean((pred - y.test)^2)
}
plot(val.errors, xlab = "Number of predictors", ylab = "Test MSE", pch = 19, type = "b")
```



Number of predictors

#(e)

```
min<-which.min(val.errors2)
# model with 20 variables has the smallest test MSE
```

(f)

```
coef(regfit.full12, min)

## (Intercept)          x.1          x.2          x.3          x.4
##  0.004523898  0.064125608  0.034223659  0.059885687  0.303085596
##          x.5          x.6          x.7          x.8          x.9
## -0.092278402  0.138192663 -0.001258715 -0.180768763 -0.175994815
##          x.10         x.11         x.12         x.13         x.14
## -0.006149230  0.194322808 -0.133014447  0.068342992  0.279067304
##          x.15         x.16         x.17         x.18         x.19
## -0.208805461 -0.345800849 -0.008159752  0.411896627  0.153556682
##          x.20
## -0.278170963
```

```
#The best model caught all zeroed out coefficients
```

(g)